

SECOND BRAIN OS

Persistent memory. Real context. AI that knows your work, your people, and your priorities.

TL;DR Five layers, two daily processes, seven skills. A vault remembers what you decide and who matters. Scheduled tasks gather and synthesize. Interactive skills decide. The result: AI that actually has your context, with a daily motion that compounds across sessions.

1 PERSISTENT KNOWLEDGE

What the system knows, and remembers between sessions

PARA Vault

Projects, Areas, Resources, Archive. Active work separated from reference. Every note backlinked. Grows with you.

Nothing is lost, everything is findable

Active Priorities

Manual strategic list. This Week / Ongoing / Watching. **Authoritative**. Morning Brief reads this first.

Strategic focus survives noisy days

Daily Notes

Morning Brief + Evening Wrap. Carry-forward loop: each evening feeds the next morning.

No day starts from zero

Memory System

People profiles, project context, working-style preferences. Makes Claude understand your org's language and shorthand.

Claude understands subtext

TASKS.md

Tactical next-actions. Maps up to Active Priorities. Bridges "what matters" and "what to do about it."

Decisions become tracked work

Why plain text. Markdown is the most AI-accessible format. The agent reads and writes files directly, with zero infrastructure between model and memory. No vendor lock-in. Works with or without AI.

2 SCHEDULED BACKBONE

Runs without you. Gathers, synthesizes, writes the daily note before you sit down.

Morning Brief default 6:45 AM

Reads Active Priorities → Calendar → Slack → Gmail (3-tier) → Drive → Notion → Asana. Writes two altitudes: **Strategic Landscape** (focus / stalled / noise) + **Ops Sweep** (respond / delegate / kill). Includes rotating Daily Challenge.

You sit down to a briefing, not an inbox

Evening Wrap default 5:05 PM

Reads today's daily note and TASKS.md. Tracks what got done, what slipped, what's still open. Writes **Leverage Assessment**, **Carry Forward** (tagged handoff), **Tomorrow's One Thing**, optional weekly **Thread**.

Honest accountability + clean handoff

The separation principle. Scheduled tasks gather and prepare. They never decide. They surface what needs deciding. The system works even when you don't open a session.

3 INTERACTIVE SKILLS

Human-in-the-loop. Decisions stay with you.

Personalize

SETUP

One-time interview. Picks your motion (strategic / coordinator / executor) from a vignette. Generates Memory.md and CLAUDE.md, then fills in the variables in your scheduled-task templates.

The package becomes your own in 15 minutes

Triggers: "/personalize-second-brain" · "set up my second brain"

Session Context

Reads Morning Brief → presents state → routes to skills. Adds delta: what changed since the briefing.

60-second context resume, any time

Triggers: "catch me up" · "where was I" · "what's on my plate" · new session start

Tactical Tasks

TASKS.md ↔ Active Priorities. Flags orphaned tasks, unpopulated priorities, stale items. Extracts commitments post-meeting.

Intent becomes tracked execution

Triggers: "update tasks" · "what's overdue" · "prioritize my week" · post-meeting

Vault Capture

Quick capture, session wrap-up, or suggested. *"That three-bucket model is worth a note. Want me to save it?"*

Nothing valuable stays in a conversation

Triggers: "capture this" · "save to vault" · "what should I capture" · proactive when frameworks emerge

Vault Cleanup

Weekly: dedup person notes (fuzzy match), fix orphaned backlinks, standardize naming, archive stale meetings, reconcile priorities ↔ tasks.

The vault stays clean enough to trust

Triggers: "clean up the vault" · "find duplicates" · suggested weekly if 7+ days since last run

Sanity Check

Audits claims against vault + external data. Rates each: confirmed, plausible, stale, contradicted, or unvalidated.

The system challenges you

Triggers: "sanity check" · "am I wrong about" · proactive when assumptions are 2+ weeks stale

Thinking Partner

Mental-model engine. First principles, inversion, pre-mortem, second-order thinking. Each matched to the situation. Detects exploring vs. defending.

Structured reasoning when stakes are high

Triggers: "help me think through" · "what am I missing" · naming any model · appears stuck

Thinking Orchestrator

PATTERN

Stakeholder-lens orchestration. Brings perspectives you wouldn't naturally apply. Anti-reinforcement guards. *Documented in customization-guide; build your own using your stakeholders.*

Pressure-tested from every angle

Status: reference pattern, not yet shipped. Extension point for personalization.

4 AGENT + INTERFACE

The model, the connectors, and how you interact with it

Claude

Opus / Sonnet

Full vault read/write. Opus for strategy + synthesis. Sonnet for speed + batch processing.

MCP Connectors

Model Context Protocol. Structured API access to external systems.

Claude Code CLI

Primary interface. Mounted to vault via Obsidian terminal. Where most work happens.

Cowork

Web/desktop interface. Slides, docs, dashboards. Same skills, different surface.

Krisp
Required · transcripts

Slack
Optional · messages

Gmail
Optional · email

Google Cal
Optional · events

Asana
Optional · tasks

Google Drive
Optional · docs

Notion
Optional · wiki

Web + APIs
Optional · research

DAILY CONTEXT CYCLE
Scheduled tasks gather. Interactive skills decide. Each day feeds the next.

6:45 AM

Morning Brief → writes today's daily note **SCHEDULED**

you open a session

on demand

Session Context → reads brief, presents state **INTERACTIVE**

as needed

Work: tasks, capture, thinking, sanity-check **INTERACTIVE**

day winds down

5:05 PM

Evening Wrap → reflection + write-back to TASKS.md **SCHEDULED**

Morning Brief reads yesterday's carry-forward. Evening Wrap creates tomorrow's. The loop compounds.

THINKING & VERIFICATION PIPELINE
Pressure-tests your claims and your thinking. No echo chambers.

01

Frame

Thinking Partner selects the right mental model for the situation

→

02

Challenge

Optional Orchestrator pattern adds stakeholder lenses + anti-reinforcement

→

03

Verify

Sanity Check audits claims against vault + external evidence

→

04

Capture

Decisions, assumptions, open questions → vault for future reference

Morning Brief Daily Challenge applies one model per day on rotation. The interactive pipeline goes deeper on demand.

This is the individual layer. One person, one vault, one agent. At enterprise scale: shared knowledge layers, role-based access, agent audit trails, managed deployment. The individual vault becomes a node in a larger network. Architecture is solved. Multi-person orchestration with governance is the unsolved problem. **Install:** [README](#) · **Setup:** [Setup Guide](#)